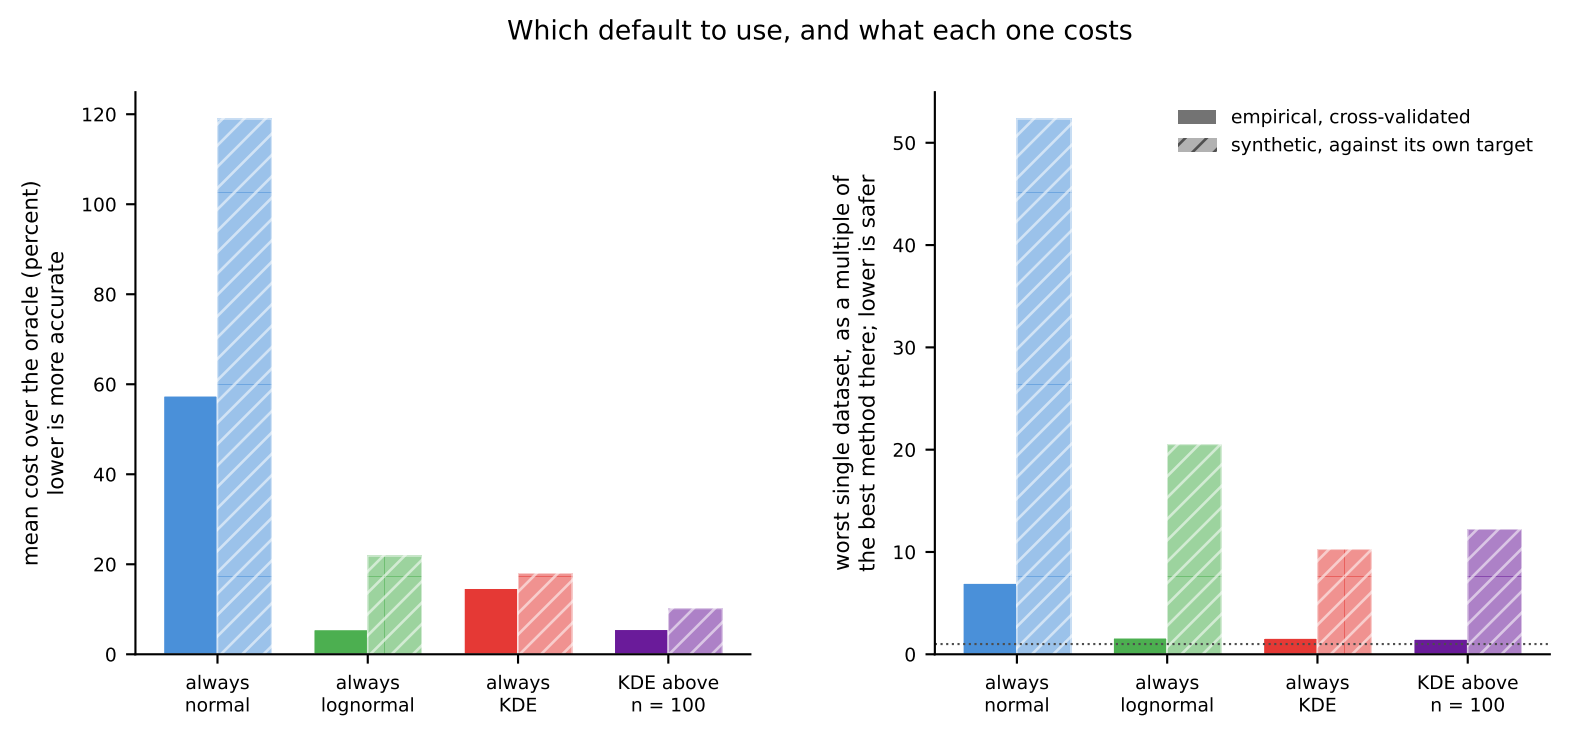
Which default to use, and what each one costs
mean cost over the oracle (percent)
lower is more accurate
0
20
40
60
80
100
120
always normal
always lognormal
always KDE
KDE above n = 100
worst single dataset, as a multiple of the best method there; lower is safer
0
10
20
30
40
50
empirical, cross-validated
synthetic, against its own target
always normal
always lognormal
always KDE
KDE above n = 100